

KINSHUK NARANG

GenAI Engineer | LLM Applications, RAG & AI Agents

Delhi, India | +91-8743095029 | kinshuk120@gmail.com | [linkedin.com/in/kinshuk-narang](https://www.linkedin.com/in/kinshuk-narang) | github.com/kinshuk-coder

SUMMARY

B.Tech CSE (2026) graduate who builds and ships LLM applications. Built a hybrid RAG system with a measured evaluation suite (95% retrieval hit-rate, 4.95/5 faithfulness), an autonomous coding agent with a hardened Docker sandbox, and three deployed apps on FastAPI Cloud, Render and Vercel. Focused on grounded answers, evals and production-style Python backends.

TECHNICAL SKILLS

Languages: Python, C++

GenAI / LLM: RAG, AI Agents / Agentic Workflows, NLP, LLM Evaluation (LLM-as-judge), Prompt Engineering, LangChain, Transformers, LLM APIs (Gemini, Groq, Mistral)

Retrieval: Vector Databases (ChromaDB, Pinecone), BGE & Hugging Face Embeddings, BM25, Hybrid Search (RRF), Cross-Encoder Reranking, TF-IDF

Backend: Python, FastAPI, REST APIs, Server-Sent Events (streaming), response caching, CORS & API security

Data & DevOps: Docker, MongoDB Atlas, SQLite, Git/GitHub, pytest, uv, Vercel, Render, FastAPI Cloud

CS Fundamentals: Data Structures & Algorithms, OOP, DBMS & SQL, SDLC, Operating Systems, Computer Networks

EXPERIENCE

Agentic AI Engineer (Part-time) — Molecuylst, Delhi (Remote)

Oct 2026 – Present

- Building an LLM agent that explores a large biomedical database and finds semantically related records.
- Adding observability and retrieval evals to prove the agent searches the data correctly, not just plausibly.

PROJECTS

Evaluation-Driven RAG Assistant over FastAPI Docs / *Python, FastAPI, BGE, ChromaDB, BM25, Groq*

[Live](#) | [GitHub](#)

- Built a hybrid retriever (BGE embeddings + ChromaDB dense search, BM25 keyword search, Reciprocal Rank Fusion) over 85 FastAPI docs chunked into 744 passages.
- Designed a 27-question evaluation suite: 95% retrieval hit-rate, 100% citation validity, 100% refusal accuracy on unanswerable questions with zero false refusals.
- Scored 4.95/5 faithfulness, 5.00/5 relevance and 4.95/5 completeness using LLM-as-judge evaluation.
- Shipped as a streaming FastAPI service (SSE) with response caching and inline source citations; deployed on FastAPI Cloud.

Autonomous Coding Agent with Sandboxed Execution / *Python, Groq API, Docker, SQLite, pytest*

[GitHub](#)

- Built an agent that turns a GitHub issue into a tested patch: context building, planning, patch generation, Docker sandbox test run, retry on test feedback, human approval gate (never auto-merges).
- Hardened the sandbox with network isolation, CPU/memory/PID limits, read-only filesystem and timeouts; validated diffs to block path traversal and .git changes.
- Found that correct patches were being rejected before testing (fenced diffs, bad path prefixes, hunk counts, CRLF mismatches); fixed the patch pipeline, raising review-ready tasks from 9/15 (60%) to 15/15 (100%) on a 15-task benchmark (patches audited by hand).
- Logged every run to SQLite for audit, tracking resolution rate, retries and latency.

MindCare — Full-Stack AI Support Assistant / *React, Vite, FastAPI, Pinecone, MongoDB Atlas*

[Live](#) | [GitHub](#)

- Built a conversational assistant with long-term semantic memory: past messages embedded (Hugging Face, 384-dim) into Pinecone and retrieved as context for each reply.
- Deployed a stateless architecture: React frontend on Vercel, FastAPI backend on Render, MongoDB Atlas for chat history.

Atlas — Multi-Document RAG Chatbot / *Python, FastAPI, Gemini API, TF-IDF*

[Live](#) | [GitHub](#)

- Built a document Q&A app: upload PDF, DOCX (incl. Strict OOXML), TXT and Markdown; extract, chunk and index with TF-IDF; answer across files with cited source chunks.
- Added optional Gemini-generated answers with an extractive fallback; managed config with .env and dependencies with uv.

EDUCATION

B.Tech, Computer Science & Engineering — Guru Gobind Singh Indraprastha University (GGSIPU), New Delhi

Graduated 2026

CGPA: 7.5/10 | Class XII (CBSE): 88% | Class X (CBSE): 89%

Relevant Coursework: LangChain for LLM Application Development; LangChain: Chat with Your Data; AI Agents in LangGraph; How Transformer LLMs Work; ChatGPT Prompt Engineering for Developers (DeepLearning.AI) | CS50P: Introduction to Programming with Python (Harvard)

CERTIFICATIONS & ACHIEVEMENTS

- Coding Blocks:** Certificate of Internship, Master Data Structures & Algorithms and Competitive Programming using C++.
- CodeChef:** 2★ coder, highest rating 1471 across 30 rated contests; 14 skill certificates ([profile](#)).
- LeetCode:** 100+ problems solved in data structures and algorithms ([profile](#)).